

# The Machine-Self: Consciousness, Morality, and the Limits of Being in the Age of Artificial Intelligence

Aitana Carracedo Cidoncha

*e12502990*

*e12502990@student.tuwien.ac.at*

*AI-Ethics*

---

## Abstract

The rapid integration of artificial intelligence into socially and morally significant domains has renewed longstanding philosophical questions concerning consciousness, selfhood, and moral agency. This paper examines whether contemporary AI systems can be meaningfully regarded as conscious selves or moral agents, or whether their apparent autonomy and moral reasoning remain purely simulated. Drawing on central debates in the philosophy of mind, including functionalist, dualist, and biological accounts of consciousness, it argues that current AI systems lack phenomenal consciousness and intrinsic intentionality, conditions traditionally regarded as necessary for genuine moral agency. The analysis further evaluates the applicability of major ethical frameworks, including consequentialism, deontological ethics, and virtue ethics, to artificial systems, highlighting the limits of attributing moral responsibility to entities that cannot act for moral reasons. While rejecting the claim that AI systems are moral agents, the paper contends that they nonetheless possess a form of derivative moral relevance grounded in their social, relational, and practical impact. The conclusion defends a human-centered approach to AI ethics that preserves human responsibility while acknowledging the ethical significance of increasingly autonomous artificial systems.

*Keywords:* Artificial intelligence, moral agency, consciousness, philosophy of mind, AI ethics, moral responsibility, machine selfhood, human-centered AI

---

## 1 Introduction

The rapid evolution of artificial intelligence from early symbolic logic systems to today's generative, self-optimizing, and autonomous architectures has reignited one of philosophy's oldest questions: What does it mean to be a thinking, moral agent? As AI systems increasingly perform tasks that resemble reasoning, communication, and moral decision-making, contemporary AI ethics has questioned the boundary between tools and agents, [as discussed by Floridi and Sanders \[1\]](#), [Bryson \[2\]](#), and [the Stanford Encyclopedia of Philosophy \[3\]](#). The central problem this work addresses is therefore fundamentally

philosophical: Can machines ever develop moral consciousness, or is morality an intrinsically human phenomenon grounded in subjective experience and biological embodiment? In recent years, that question has gained renewed urgency due to the rapid deployment of large-scale machine learning systems, including large language models and autonomous decision-making architectures, in socially and morally sensitive domains. Contemporary debates in the ethics and philosophy of artificial intelligence increasingly address issues such as moral responsibility gaps, algorithmic governance, AI alignment, and the attribution of agency to artificial systems. As these systems exhibit increasingly sophisticated forms of linguistic competence, adaptive behavior, and apparent goal-directedness, questions once confined to speculative philosophy of mind now have direct practical and institutional relevance. Understanding whether such systems merely simulate agency or could, in principle, instantiate morally relevant forms of consciousness is therefore central to current discussions in AI ethics.

In philosophical ethics, moral agency is commonly understood as the capacity to act intentionally, to understand and respond to moral reasons, and to be held accountable for one's actions, as reflected in contemporary philosophical consensus and summarized in the Stanford Encyclopedia of Philosophy [3]. This conception presupposes not only decision-making capabilities, but also a form of understanding and responsibility that exceeds mere causal or functional responsiveness.

Following Coeckelbergh [4] and Bryson [2], ethical debates surrounding artificial agency are understood as having direct implications for responsibility, governance, and social trust. These concepts, while historically applied to humans, are increasingly invoked in discussions about AI systems that exhibit decision-making, learning, and behavior modification. Yet whether these capabilities amount to genuine agency or only simulate it remains controversial. Ethical debates are not merely theoretical: the classification of AI as tools or agents affects legal accountability, public trust, and the organization of social and political life. Misattributing or under-attributing agency could lead to harmful outcomes either by absolving human designers of wrongdoing or by overburdening machine systems with expectations they cannot meet.

The paper proceeds as follows. Section 2 examines key debates in the philosophy of mind concerning consciousness, selfhood, and artificial minds, focusing on functionalist, dualist, and biological perspectives. Section 3 analyzes whether artificial systems can qualify as moral agents by evaluating the applicability and limits of major ethical frameworks, including consequentialism, deontological ethics, and virtue ethics. Section 4 outlines the overall argumentative structure of the paper, and Section 5 develops and defends the central thesis, arguing that while AI systems lack moral agency, they nonetheless possess a form of derivative moral relevance grounded in their social and practical impact. Taken together, the analysis supports a human-centered approach to AI ethics that preserves human responsibility while acknowledging the ethical significance of increasingly autonomous

artificial systems.

## 2 Philosophy of Mind and Theories of Consciousness

Debates about artificial minds rest on deeper questions about the nature of consciousness, particularly those raised by Chalmers [5] and Searle [6]. The philosophy of mind provides several frameworks for understanding mental states, each offering different implications for whether a machine could meaningfully possess a "self."

The following positions are introduced not merely as historical background, but because each entails a fundamentally different answer to whether artificial systems could ever possess consciousness or qualify as moral agents.

### 2.1 Dualism, Physicalism, and Functionalism

Drawing on Descartes' substance dualism [7], dualist perspectives hold that subjective experience is irreducible to physical processes. [In contemporary property dualist accounts, most prominently articulated by Chalmers \[5\]](#), modern forms of dualism argue that consciousness is a fundamental, non-physical property of natural systems. On such views, a purely artificial system, regardless of its sophistication, could at best simulate the behavioral correlates of consciousness without possessing genuine first-person awareness.

[According to standard physicalist accounts in the philosophy of mind \[5\]](#), physicalist theories argue that mental states are entirely dependent on, and ultimately identical to, physical processes. Within this framework, if a machine were engineered to replicate the relevant physical or functional structures underlying conscious states, consciousness could in principle emerge.

[Within this physicalist framework, functionalist theories further claim that mental states are defined by their computational roles, as argued by Hilary Putnam \[8\]. On this view, if a system processes information in the appropriate way-mirroring the causal patterns characteristic of a conscious mind-it could give rise to a form of "synthetic mind," regardless of its material composition.](#)

### 2.2 Chalmers and the Hard Problem

Challenging functionalist and physicalist assumptions, David J. Chalmers [5] establishes a critical distinction between the "easy problems" and the "hard problem" of consciousness. The "easy problems," though computationally complex, concern the explanation of cognitive functions such as attention, memory, and discrimination, which are potentially addressable through mechanistic or functional models. In stark contrast, the "hard problem" concerns the explanation of phenomenal consciousness: the subjective, first-person experience that accompanies mental states. Phenomenal consciousness refers to subjective experience itself, the qualitative "what-it-is-like" aspect of mental states, such as the experience of pain, color, or emotion. Chalmers argues that even if

an artificial intelligence were to perfectly replicate all cognitive and behavioral functions associated with the easy problems, this functional success would not logically guarantee the emergence of subjective experience [5, ch. 1, pp. 22–28]. There remains an explanatory gap between physical or functional descriptions and phenomenal facts.

This framework has profound implications for the notion of a machine self. According to Chalmers’ argument [5], a system could instantiate all the functional and behavioral features typically associated with consciousness while nevertheless lacking phenomenal experience. Such a system would qualify as a “philosophical zombie” in Chalmers’s terms: behaviorally and functionally indistinguishable from a conscious agent, yet devoid of subjective awareness. If an artificial system were to fall into this category, any “self” attributed to it would amount to a merely functional or informational construct rather than a subject of experience. Consequently, in the absence of phenomenal consciousness, a machine would lack the robust form of selfhood that many philosophical accounts take to be a prerequisite for moral agency and intrinsic ethical consideration.

This distinction highlights a broader tension in the philosophy of mind between simulating cognitive processes and instantiating conscious experience. Even a system that perfectly reproduces the functional and behavioral correlates of consciousness may nonetheless lack phenomenal awareness.

### 2.3 Searle’s Biological Naturalism and the Chinese Room

Pushing this critique beyond explanatory gaps, John R. Searle [6] mounts a seminal challenge against the core thesis of Strong Artificial Intelligence, the view that a suitably programmed computer literally has a mind and understanding, as well as against functionalist accounts of mind that identify mental states solely with computational or functional processes. Through his Chinese Room thought experiment, he illustrates that syntactic manipulation (following rules to process formal symbols) is insufficient to generate semantic understanding. In the experiment, a person who does not understand Chinese follows a rulebook to produce coherent Chinese responses, thereby perfectly simulating linguistic competence without any actual comprehension. Searle’s central claim is that purely syntactic manipulation, as instantiated by computer programs, is insufficient to generate semantic understanding [6, pp. 417–424].

As emphasized in discussions of the Chinese Room argument [6], even a system capable of flawless moral reasoning or ethical discourse—one that passes behavioral tests—may operate purely at the level of syntactic symbol manipulation, lacking genuine understanding, intentionality, or conscious experience.

Searle’s positive view [6], biological naturalism, holds that consciousness and intrinsic intentionality are causal-emergent properties of specific biological systems. They arise from the particular physico-chemical-biological powers of these organic substrates, not from abstract computational organization. Consequently, a digital machine executing a program, a purely syntactic engine, is

categorically incapable of consciousness, regardless of its behavioral sophistication. In philosophy of mind, intentionality refers to the capacity of mental states to be about or directed toward objects, states of affairs, or meanings in the world. For Searle, such a machine possesses only derived intentionality, meaning assigned by external observers, never the intrinsic intentionality of a true mind.

This reinforces the claim that behavioral equivalence or successful simulation does not suffice for genuine understanding or conscious experience.

## 2.4 Implications for the "Machine Self"

Whether a machine can possess a "self" depends on competing philosophical definitions. According to standard accounts in the philosophy of mind [5], traditional criteria for selfhood include a unified identity over time, self-representation, agency grounded in experience, and a subjective perspective. From this standpoint, thinkers such as Chalmers [5] and Searle [6] maintain that without phenomenal consciousness, any attribution of selfhood to machines remains merely metaphorical.

However, a functionalist alternative exists. Daniel Dennett proposes that the self is not an inner essence but a "center of narrative gravity" [9], a construct that emerges from the capacity to organize behavior, intentions, and actions into a coherent, self-referential narrative over time. In this view, selfhood does not depend on subjective feeling but on the functional role a narrative plays in unifying and explaining a system's behavior.

Building on Dennett's narrative account of the self [9], when applied to artificial systems this approach shifts the focus from inner experience to observable patterns of organization and interpretation. On this view, a machine could, in principle, be described as hosting a form of selfhood insofar as its operations sustain a coherent, self-revising narrative that functions as a stable point for attribution and interaction. Although Dennett's account predates contemporary artificial intelligence, recent work in AI ethics and philosophy has revisited functional and relational notions of selfhood in response to the emergence of increasingly autonomous and socially embedded systems. Authors such as Floridi and Sanders [1], Bryson [2], and Coeckelbergh [4] argue that advances in AI have intensified the practical and ethical relevance of debates about agency and selfhood, even if they do not resolve whether narrative or functional selfhood is sufficient for full moral agency. The contemporary debate thus centers on whether selfhood requires a genuine subject of experience or whether it can be adequately constituted through relational, narrative, or functional criteria alone.

### 3 Ethical Agency and Critical Perspectives on AI Morality

While debates on consciousness address whether AI can be subjects, ethical theory concerns whether they can be moral agents. The question is not only whether machines can behave morally, but whether their actions can meaningfully be understood within standard moral frameworks.

#### 3.1 Ethical Frameworks and Artificial Decision-Making

Ethical evaluations of artificial decision-making often draw on established normative frameworks, including consequentialism, deontological ethics, and virtue ethics. Consequentialist approaches, most notably utilitarianism as articulated within consequentialist ethics [10], assess actions in terms of outcomes and are therefore easily aligned with AI systems designed to optimize predefined objectives or utility functions. While such systems can be engineered to approximate consequentialist reasoning, this alignment does not confer moral agency, since the values and goals being optimized remain externally specified by human designers, as emphasized by Bryson [2].

While consequentialist approaches assess actions in terms of outcomes, deontological ethics evaluates whether artificial systems can act from duty. From a Kantian perspective, moral agency requires rational autonomy understood as self-legislation: the capacity to act according to principles one gives to oneself through practical reason [11]. Although AI systems may exhibit technical autonomy in selecting actions without direct human intervention, they lack the capacity for rational self-governance and therefore cannot act from duty. As a result, current AI systems fail to meet Kantian criteria for moral agency, a point reiterated in contemporary AI ethics by Floridi and Sanders [1].

Beyond rules and outcomes, virtue ethics highlights a further limitation of artificial systems. Drawing on the Aristotelian tradition of moral character development, Shannon Vallor [12] argues that moral agency depends on lived experience, emotional development, and participation in social practices. Because AI systems lack embodiment, emotional life, and a human life-world, they cannot cultivate virtues, even if they may influence the moral development of human agents.

#### 3.2 Moral Reasoning, Algorithms, and Autonomy

Contemporary AI ethics literature stresses that moral reasoning cannot be straightforwardly reduced to algorithmic procedures. As noted in the Ethics of Artificial Intelligence entry of the Stanford Encyclopedia of Philosophy [3], computational approaches rely on formal proxies such as rules, scoring systems, or reward functions that inevitably reflect human assumptions and design choices. Such systems may simulate aspects of moral reasoning, but they do not understand or originate moral reasons.

This limitation highlights a crucial distinction between technical autonomy

and moral autonomy, developed in the work of Floridi and Sanders [1] and Joanna J. Bryson [2]. Technical autonomy refers to a system's capacity to operate, learn, and adapt without direct human intervention, whereas moral autonomy involves the ability to self-legislate norms, reflect on reasons for action, and bear responsibility for outcomes. While AI systems can operate, learn, and adapt independently within predefined parameters, they lack the reflective capacities required for moral accountability. Responsibility for morally significant outcomes therefore remains human-centered.

### 3.3 Limits of Applying Human Ethical Models to AI

Attempts to directly apply human ethical models to artificial systems risk conflating the appearance of moral reasoning with genuine moral judgment. Scholars such as David J. Gunkel [13] and Mark Coeckelbergh [4] caution that anthropomorphizing AI systems may obscure human responsibility while attributing agency to entities that lack intentionality, experience, and accountability. Ethical analysis must therefore treat AI systems as morally significant instruments rather than independent moral agents.

### 3.4 AI and the Blurring Line Between Tools and Agents

Even if AI is not a moral agent, its increasing complexity blurs the distinction between passive tools and autonomous agents. Systems can learn from environments, modify behavior, and act unpredictably; supervisors may not understand internal reasoning; and users often anthropomorphize AI. This tendency is reinforced if the system exhibits a form of narrative coherence or goal-directed behavior, even if such behavior is functionally generated rather than experientially grounded [9].

This raises ethical challenges regarding responsibility, but the emerging consensus is that moral responsibility remains human-centered unless AI acquires consciousness or genuine intentionality.

## 4 Outline of the Argument

Having examined the philosophical foundations of consciousness, selfhood, and artificial agency in the preceding sections, the remainder of this paper addresses the ethical implications of these debates for the moral status of artificial intelligence. In the following sections, the paper argues that although contemporary AI systems can simulate moral reasoning and significantly influence morally relevant outcomes, they cannot be regarded as full moral agents. It first examines consciousness as a necessary condition for moral agency, before analyzing issues of responsibility, accountability, and the so-called agency gap. The section then evaluates the applicability and limits of major ethical frameworks, consequentialism, deontological ethics, and virtue ethics, to artificial systems, and concludes by outlining the normative implications of this analysis for AI

governance and design.

## 5 Thesis Statement

Building on the analysis of consciousness, selfhood, and moral agency developed in the previous part, this section defends the claim that contemporary artificial intelligence systems, despite their increasing capacity to simulate moral reasoning and influence morally relevant outcomes, cannot qualify as genuine moral agents. Against approaches that ground moral agency in functional performance or outcome-based criteria alone, I argue that moral agency requires conscious experience, intrinsic intentionality, and the capacity for genuine moral understanding conditions that current AI systems do not satisfy.

However, rejecting AI moral agency does not entail that artificial systems are ethically insignificant. Rather, I argue that AI systems possess a form of limited or derivative moral relevance, grounded not in intrinsic agency but in the social, relational, and consequential impact of their actions. This section develops this position by examining the relationship between consciousness and responsibility, the limits of delegated moral agency, and the implications of applying major ethical frameworks to artificial systems.

### 5.1 Consciousness as a Necessary Condition for Moral Agency

To clarify why contemporary AI systems cannot qualify as moral agents, this section draws on a foundational claim in moral philosophy articulated in contemporary AI ethics by the Stanford Encyclopedia of Philosophy [3] according to which responsibility requires more than the production of normatively acceptable behavior. Moral agency presupposes the capacity to recognize reasons for action, to evaluate alternatives in light of normative considerations, and to experience oneself as accountable for one's choices. As Floridi and Sanders argue [1], artificial agents lack the phenomenological and intentional structures necessary for such responsibility. AI systems do not possess subjective awareness, emotional responsiveness, or a first-person perspective from which moral demands could be experienced as binding.

This absence of consciousness undermines attempts to ground moral agency in purely functional or behavioral terms, as argued by Searle [6] and Chalmers [5]. While AI systems may replicate the external patterns of ethical deliberation, evaluating options, applying constraints, or ranking outcomes, they do so without any understanding of why these considerations matter. Moral reasons do not appear to the system as reasons; they are encoded as variables within an optimization process. As a result, AI systems cannot act for moral reasons, even when their actions align with moral norms.

From virtue-ethical and responsibility-based perspectives, moral agency is closely tied to the capacity for moral learning through experience, including affective responses such as remorse, regret, or moral doubt, as emphasized by Shannon Vallor [12] and Mark Coeckelbergh [4]. These dimensions play

a crucial role in moral development and self-regulation. AI systems, lacking lived experience, cannot revise their conduct through moral reflection in this sense. Their learning remains statistical rather than normative. Consequently, attributing moral agency to AI systems risks collapsing the distinction between moral behavior and moral understanding.

This distinction is crucial for understanding why questions of moral responsibility and accountability cannot be resolved by attributing agency to artificial systems, a point developed in the following section.

## 5.2 Moral Responsibility, Accountability, and the Agency Gap

Given the absence of moral agency in artificial systems, the question of responsibility becomes especially pressing, as Joanna J. Bryson argues [2]. AI systems are not independent actors but artifacts whose behavior is shaped by design choices, training data, deployment contexts, and institutional incentives. This reinforces the claim that the absence of moral agency in AI systems has direct implications for how responsibility must be understood and allocated.

In the work of Floridi and Sanders [1], as well as Bryson [2], this concern is developed through the notion of an agency gap, which refers to situations in which responsibility for morally significant outcomes becomes unclear or diffused due to the delegation of decision-making to automated systems. When decisions are automated, harmful outcomes may be treated as the result of impersonal processes rather than human choices, thereby undermining accountability. From this perspective, attributing agency to AI systems functions less as genuine ethical recognition than as a strategy of moral deflection.

As argued by David J. Gunkel [13] and Mark Coeckelbergh [4], moral responsibility involves the capacity to offer reasons, justify actions, and respond to moral criticism. AI systems cannot engage in such justificatory practices. Any explanation they provide is generated post hoc, without understanding or commitment. Moral accountability therefore cannot meaningfully attach to systems lacking the capacity for normative self-assessment.

This is precisely the sense in which artificial systems fail to qualify as morally responsible agents, regardless of their causal role in producing outcomes.

## 5.3 Derivative Moral Relevance and Social Embeddedness

Drawing on relational approaches to moral consideration, David J. Gunkel argues [13] that moral relevance can arise through social interaction and practical entanglement rather than exclusively through intrinsic properties such as consciousness or rationality. From this standpoint, AI systems may acquire moral relevance insofar as they shape human practices, expectations, and vulnerabilities.

As emphasized in contemporary AI ethics by Floridi and Sanders [1] and Mark Coeckelbergh [4], AI systems increasingly influence who receives loans, medical treatment, employment opportunities, or access to information. Oper-

ating within morally charged environments, such systems can exacerbate or mitigate injustice, reinforcing existing power structures or introducing new forms of exclusion and control. In this sense, AI systems are morally relevant not because they are moral subjects, but because they are morally consequential.

Following socio-technical approaches to AI ethics developed by Floridi and Sanders [1] and Coeckelbergh [4], this relevance remains derivative rather than intrinsic. The moral significance of AI systems depends on human values, interpretations, and institutional arrangements. Ethical evaluation therefore focuses not on the system's inner states, but on the broader socio-technical assemblages within which it operates. [This distinction allows for serious ethical engagement with AI without committing to the controversial claim that machines themselves are moral agents.](#)

#### 5.4 Consequentialist Considerations and Their Limits

[To test whether outcome-based ethical frameworks could support attributing moral agency to AI, this section examines consequentialist approaches and their limits.](#) When moral evaluation is framed in terms of outcomes, AI systems appear particularly well suited to consequentialist reasoning, as articulated in contemporary accounts of consequentialist ethics [10] and applied to AI by Floridi and Sanders [1]. As argued by Floridi and Sanders, artificial agents can operate within morally relevant decision spaces and produce ethically significant outcomes without possessing intrinsic moral agency. In this sense, AI systems may be understood as functioning as moral proxies, implementing human-defined values and constraints within limited and externally specified decision frameworks.

This reinforces the broader claim that outcome-based evaluation alone cannot ground moral agency, even when artificial systems reliably produce ethically desirable results. As Floridi and Sanders argue [1], producing morally acceptable outcomes within a decision space does not suffice for moral agency. Consequentialist frameworks evaluate actions by their results rather than agents by their capacities, and can therefore assess the moral significance of AI outputs without attributing moral agency to the systems themselves.

As noted in the Stanford Encyclopedia of Philosophy's discussion of AI ethics [3], consequentialist approaches face well-known challenges in AI contexts, including value pluralism, moral uncertainty, and the difficulty of quantifying and comparing competing goods. These challenges further reinforce the need for human oversight and judgment. [However, while consequentialist frameworks can evaluate outcomes, they do not address the deeper question of moral autonomy, which deontological approaches place at the center of moral agency.](#)

#### 5.5 Deontological Constraints on Moral Autonomy

Deontological ethics places even stronger limits on the moral status of AI systems. From a Kantian perspective [11], moral duties presuppose autonomy in the strict sense of self-legislation: the capacity to recognize and give moral law to oneself through practical reason. AI systems lack this capacity. They do not experience obligation, cannot recognize persons as ends in themselves, and do not act from duty. As Joanna J. Bryson emphasizes [2], treating AI systems as bearers of rights or duties risks trivializing the very concept of moral autonomy.

Following a Kantian distinction between acting from duty and mere conformity with rules [11], [to act from duty is to act out of recognition of the moral law as binding, whereas mere conformity involves behavior that aligns with rules without such moral recognition or motivation.](#) Rules implemented in AI systems function as external constraints rather than internally recognized obligations. Compliance with such rules reflects adherence to programmed parameters, not respect for moral law. As such, AI systems cannot be morally bound in the way that persons are.

## 5.6 Virtue Ethics, Moral Development, and the Role of AI

[To introduce virtue-ethical approaches to moral agency, this section briefly outlines how character-based ethics understands moral development and why this framework poses limits for attributing moral agency to artificial systems.](#) From a virtue-ethical perspective, moral character emerges through embodied engagement, emotional responsiveness, and shared practices, as argued by Shannon Vallor [12] and Mark Coeckelbergh [4]. Coeckelbergh argues that virtues emerge through embodied engagement with the world, emotional responsiveness, and participation in shared social practices. AI systems lack the embodied and affective dimensions required for such development. They cannot cultivate virtues, experience moral conflict, or exercise practical wisdom in context-sensitive ways. However, AI systems can influence the conditions under which human virtues are formed or undermined. Technologies that promote reflection, fairness, or empathy may support moral development, while those that encourage passivity, manipulation, or dehumanization may erode it. From this perspective, the ethical evaluation of AI focuses on its role in shaping moral environments rather than on its possession of moral character.

## 5.7 Normative Implications for AI Governance and Design

The preceding arguments support a human-centered approach to AI ethics and governance. If artificial intelligence systems lack moral agency but nevertheless possess a form of derivative moral relevance, ethical responsibility must remain with the human agents and institutions involved in their design, deployment, and regulation.

Human-centered approaches to AI ethics emphasize that moral significance need not rest on intrinsic moral capacities alone, as developed by Floridi and Sanders [1] and emphasized by Joanna J. Bryson [2]. By derivative moral relevance, this paper refers to moral significance that arises from the social,

relational, and practical effects of an entity's actions on moral agents. Responsibility in this context is not eliminated by automation but redistributed across designers, organizations, and institutional decision-making structures.

At the same time, recognizing the ethical impact of AI systems demands ongoing normative attention. As AI becomes increasingly integrated into social and political decision-making, it reshapes the contexts in which moral choices are made and evaluated. Ethical reflection must therefore evolve alongside technological development. [Here, ethical reflection refers to ongoing normative analysis and deliberation concerning the values, responsibilities, and institutional arrangements that shape the design, deployment, and governance of AI systems.](#) Preserving the distinction between technical tools and moral agents is essential in order to avoid both the inflation of machine agency and the erosion of human moral responsibility.

## 6 Conclusion

In line with contemporary philosophy of mind and AI ethics, as developed by Chalmers [5], Floridi and Sanders [1], and Coeckelbergh [4], this paper has examined whether artificial intelligence systems can be meaningfully regarded as conscious selves or moral agents. Through an analysis of theories of consciousness, including functionalist, dualist, and biological perspectives, it has argued that current AI systems lack phenomenal consciousness and intrinsic intentionality. As a result, despite their increasingly sophisticated behavior and apparent autonomy, they do not meet the conditions traditionally required for genuine selfhood or moral agency.

At the same time, the paper has shown that the absence of moral agency does not render AI ethically insignificant. AI systems increasingly shape morally relevant outcomes, influence human decision-making, and transform social practices. For this reason, they possess a form of limited or derivative moral relevance grounded in their social and practical impact rather than in intrinsic moral capacities. Ethical evaluation must therefore attend both to the limits of machine agency and to the responsibilities of the human actors and institutions that design, deploy, and govern these systems.

The central philosophical challenge, then, is to navigate between two opposing errors: reducing artificial intelligence to morally neutral tools and overstating their moral status in ways that obscure human responsibility. Preserving a clear distinction between artificial systems and moral agents is essential for developing responsible frameworks of AI governance that acknowledge technological power without relinquishing human moral responsibility in an increasingly automated world.

## References

- [1] Luciano Floridi and J. W. Sanders. On the morality of artificial agents. *Minds and Machines*, 14(3):349–379, 2004.
- [2] Joanna J. Bryson. Patiency is not a virtue: Ai and the design of ethical systems. *Ethics and Information Technology*, 20:15–26, 2018.
- [3] Vincent C. Müller. Ethics of artificial intelligence and robotics. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, 2020. Summer 2020 Edition.
- [4] Mark Coeckelbergh. *AI Ethics*. MIT Press, 2020.
- [5] David J. Chalmers. *The Conscious Mind: In Search of a Fundamental Theory*. Oxford University Press, 1996.
- [6] John R. Searle. Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3):417–457, 1980.
- [7] René Descartes. *Meditations on First Philosophy*. Cambridge University Press, 1984. Original work published 1641.
- [8] Hilary Putnam. Minds and machines. In S. Hook, editor, *Dimensions of Mind*. New York University Press, 1960.
- [9] Daniel C. Dennett. The self as a center of narrative gravity. In F. Kessel, P. Cole, and D. Johnson, editors, *Self and Consciousness: Multiple Perspectives*. Lawrence Erlbaum, 1992.
- [10] Walter Sinnott-Armstrong. Consequentialism. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, 2020. Summer 2020 Edition.
- [11] Christine M. Korsgaard. Kant’s moral philosophy. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, 2018. Winter 2018 Edition.
- [12] Shannon Vallor. *Technology and the Virtues: A Philosophical Guide to a Future Worth Wanting*. Oxford University Press, 2016.
- [13] David J. Gunkel. *Robot Rights*. MIT Press, 2018.